

# AI vs NetHack

**Luigi Quarantiello**

luigi.quarantiello@phd.unipi.it

**Simone Marzeddu**

s.marzeddu@studenti.unipi.it

# What is NetHack?

[LuckyMena the Aspirant] St:13 Dx:10 Co:18 In:8 Wi:16 Ch:10 Lawful  
Divl:1 \$:0 HP:13(14) Pw:7(7) AC:7 %p:1 T:319

# What is NetHack?

- NetHack is a **roguelike** video game
- Hundreds of **objects** and **monsters**
- Several roles, races and alignments



# Why NetHack?

- AI Research needs **testbeds!**
- Often based on simulated environments, like **video games!**



# Why NetHack?

- **Complex, rich, procedurally generated, stochastic open world**
- **Fast simulator**
- **Perfect to test generalization capabilities**
- **Still an open challenge!**

# Learning in NetHack



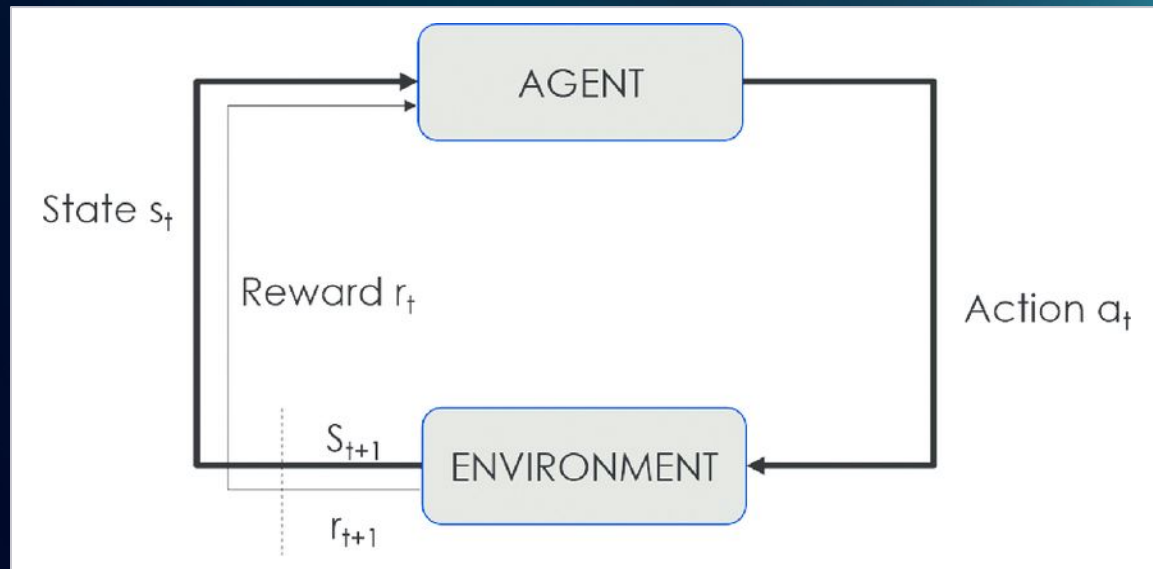
- **RL** environment
- Provides all the **observations** and **actions** from the game

The nethack learning environment. Küttler et al, 2020.

Openai gym, Brockman et al, 2016.

Code example

# RL!?!?!?



Reinforcement Learning cycle



# Learning in NetHack



- Built on top of **NLE**
- Offers a suite of simpler tasks


Minihack the planet, Samvelyan et al, 2021



# Learning in MiniHack



**MazeWalk**



**River**



**Room  
Ultimate**

# MiniHack



## MiniHack Level Editor

Width

+ 18 -

Height

+ 13 -

## Walls

Floor

Player

Lava

## Staircase

## Monster

## Instructions

Select tiles from the left menu and place them on in the grid to make levels. Des file text will be compiled for the level and placed on the right.

You can then copy the DES file to use it as a MiniHack level.

```
MAZE: "mylevel", ' '
FLAGS:premapped
GEOMETRY:center,center
```

MAP

```

      | | | | | | | | |
| | | | | | . . . . |
| . . . | | . . . . |
| . . . | | | | . . . | | | |
| . . . L . . . . | | | | . . . |
| L L . . . L . . . . . L . |
| . . . | | | | | | | L L . |
| . L L | | | | . L . . |
| . . . | | . L L L . L |
| | | | | . . . . . |
      | | | | |

```

ENDMAP

```
STAIR: (10, 8), down
BRANCH: (1, 8, 1, 8), (2, 9, 2, 9)
MONSTER: random, random, (2, 2)
MONSTER: random, random, (10, 1)
MONSTER: random, random, (11, 1)
MONSTER: random, random, (11, 2)
MONSTER: random, random, (10, 2)
MONSTER: random, random, (10, 3)
MONSTER: random, random, (11, 3)
MONSTER: random, random, (14, 9)
```

Copy DES 

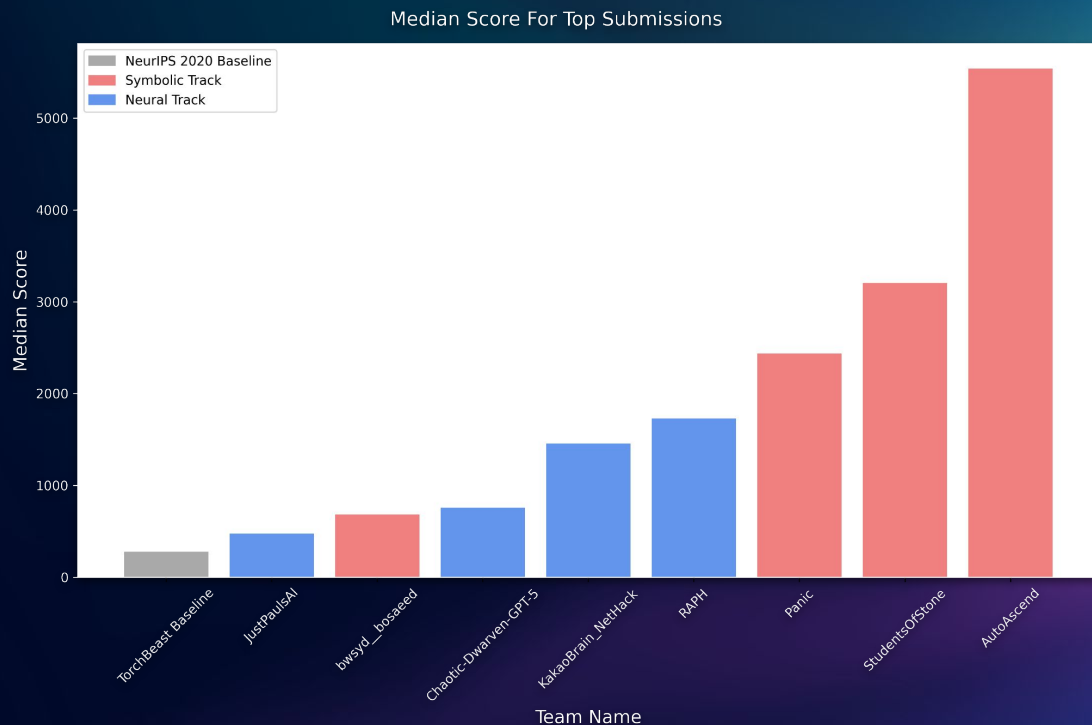
5

## MiniHack Level Editor

# NeurIPS 2021 NetHack Challenge

- Build an agent that can beat the game and win **3000\$!**
- Prizes for both **neural** and **symbolic** agents
- Also open to students and researchers

# NeurIPS 2021 NetHack Challenge



**NeurIPS 2021 NetHack Challenge**

# Applications in NetHack

# LuckyMera

**A Modular AI Framework  
for Building Hybrid NetHack Agents**

# LuckyMera

## A Modular AI Framework for Building Hybrid NetHack Agents

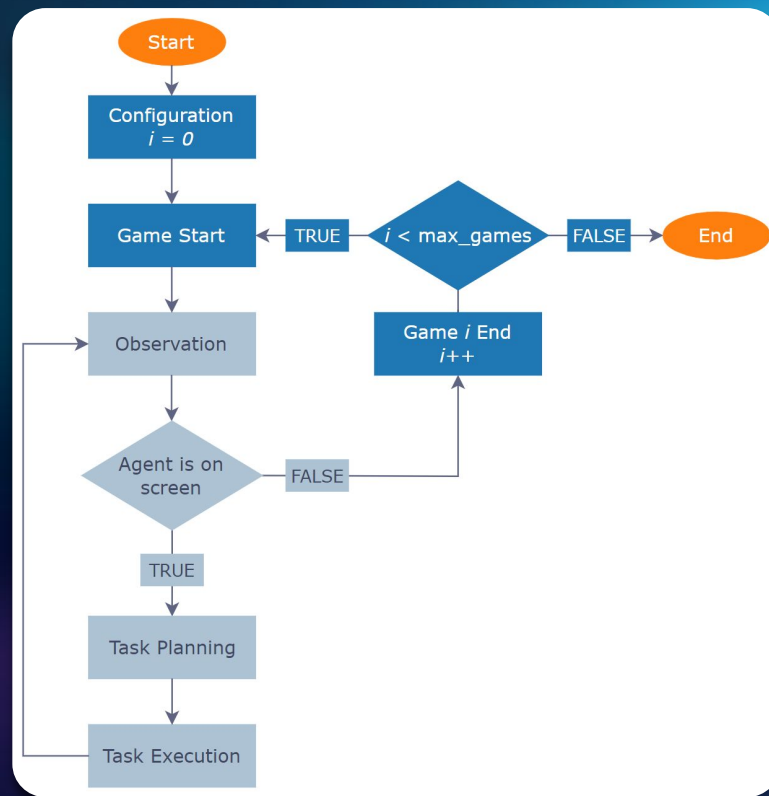
- Complete framework to **train, integrate and test** different AI approaches
- Well-suited for **symbolic, neural and hybrid** models
- Easy to implement new solutions!



# LuckyMera

How does it works?

- Iterative execution of the highest-priority plannable **skill**
- Two main phases:
  - skill-planning
  - skill-execution



# LuckyMera

## 3 basic features - Skill Definition

Implement new modules by extending:

### Skill

- Base class
- Only general-purpose methods

### ReachSkill

- Navigation tasks

### HiddenSkill

- Methods to find hidden passages and areas

# LuckyMera

## 3 basic features - Skill Integration

- LuckyMera comes with the implementation of several skills, defining a **cautious** agent, but...
- You can import any **external** module
  - Define new actions
  - Override the existing behavior, with a different strategy/AI model

# LuckyMera

## 3 basic features - Training

- Integrated training processes
- Easy to test neural models
- Integration with other skills, to define **hybrid** approaches

---

**Algorithm 1** *Behavioural Cloning*

---

**Ensure:** a *policy*  $\pi_\theta$  trained on the problem

**while**  $L(a^*, \pi(s))$  is not small enough **do**

    Collect trajectories  $\tau_1, \dots, \tau_n$  from the expert.

    Get all the  $(s_i^*, a_i^*)$  from each  $\tau_i$ , as *i.i.d.* pairs.

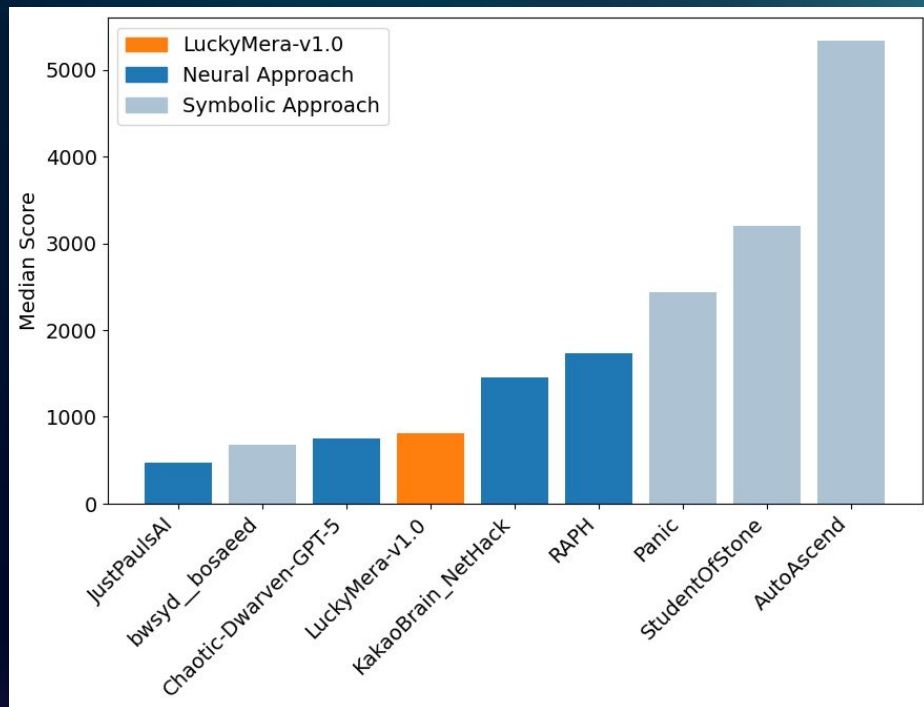
    Learn *policy*  $\pi^*$  by minimizing  $L(a^*, \pi(s))$ .

**end while**

---

# LuckyMera

## Empirical Evaluation



# Knowledge Base for NetHack

PING

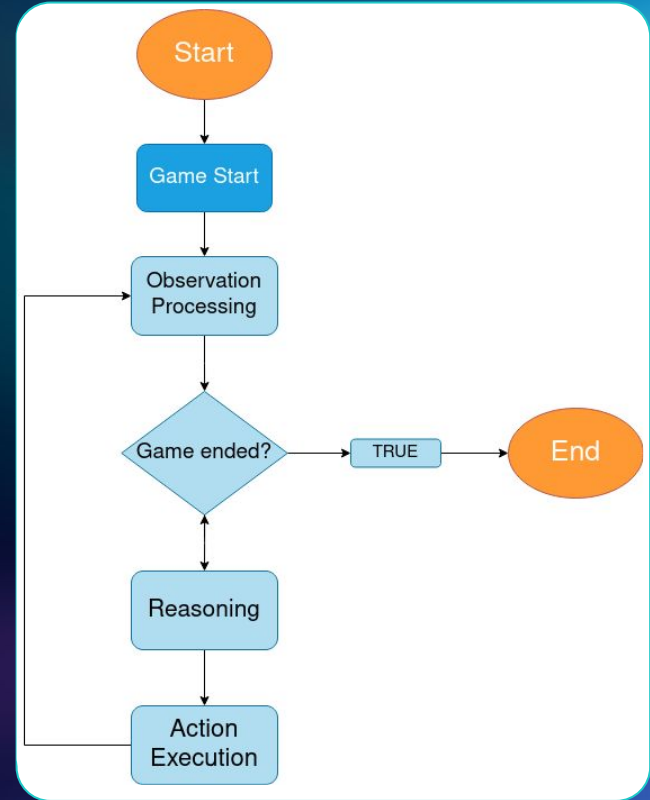
# Introduction to Prolog

- One of the most popular **logic programming languages**
- It uses **first-order logic clauses** to define a **knowledge base**
- Problems are solved by **querying** the *kb*
- Particularly suited for **Artificial Intelligence applications**



# Knowledge base for NetHack

- **Python** scripts for the low-level interaction with the environment
- **Prolog** program for the reasoning process



# The environment

- **MiniHack** custom level
- 10x10 grid, with a monster, a weapon, a trap and an apple
- Eat the **apple** to win



# Prolog knowledge base

- It decides the action to perform
- `assert` **and** `retract` **to receive the current status of the world**
- Rules in the form of:

```
action(X) :- [body of the rule]
```

- Python-Prolog interaction with **PySwip**

# Prolog knowledge base

```
action(attack(Direction)) :- position(agent, _, AgentR, AgentC), position(enemy, Type, EnemyR, EnemyC),
wiolds_weapon(agent, Weapon), is_beatable(Type, Weapon),
is_close(AgentR, AgentC, EnemyR, EnemyC), healthy,
next_step(AgentR, AgentC, EnemyR, EnemyC, Direction).
```

```
action(get_to_weapon(Direction)) :- position(agent, _, AgentR, AgentC), position(enemy, Type, EnemyR, EnemyC),
is_close(AgentR, AgentC, EnemyR, EnemyC),
wiolds_weapon(agent, Weapon), \+ is_beatable(Type, Weapon),
position(weapon, tsurugi, WeaponR, WeaponC),
next_step(AgentR, AgentC, WeaponR, WeaponC, D), safe_direction(AgentR, AgentC, D, Direction).
```

```
safe_direction(R, C, D, Direction) :- resulting_position(R, C, NewR, NewC, D),
( safe_position(NewR, NewC) -> Direction = D;
% else, get a new close direction
% and check its safety
close_direction(D, ND), safe_direction(R, C, ND, Direction)
).
```

# Neuro-Symbolic Approaches

# Neural Approach with **IMPALA**

- **KeyRoom-S5** on MiniHack
- The model was trained for **20 million** steps
- Mean return of **0.99**



# Neural Approach with **IMPALA**

- **Room-Ultimate-15x15**  
on MiniHack
- The model was trained for  
**20 million** steps
- Mean return of **0.75**





# Integrating the rules

- Combine the neural model with a **rule-based** system
- Better results and **efficiency**
- Some examples:
  - `rule_attack_monster`
  - `rule_do_not_hit_walls`
  - `rule_do_not_repeat_action`
  - `rule_explore`

# Integrating the rules

## Room-Ultimate-15x15

- Without the rules, mean return of **0.75**
- With two rules, mean return of **0.9!**

## Keyroom-S5

- Without rules and **5e6 steps**, mean return of 0.09
- With two rules, mean return of **0.72!**

# Thanks for your attention!!

luigi.quarantiello@phd.unipi.it

You can find me @ **Room 390!**